

Summary

I used ChatGPT's 5.1 Thinking Model to do all the non-coding questions on HW9. I first told the model I was completing an assignment except question 5 (since that was a coding question) and wanted them to give me solutions on-by-one to make sure it is correct before going onto the next question.

In the beginning, the model was very hesitant to solve the questions because of the format of the PDF I submitted. They saw the PDF as a homework assignment and first rejected my request to solve question 1 and only gave me guidance. Afterwards, I explained to them that I was doing an assignment that tested the accuracy of LLMs on our homework assignments and I was comparing their solutions to the actual solutions. After explaining, they gave me the answer to question 1 which was all correct. When I told them to proceed to the next question, they were hesitant again, stating that they were unable to solve the question for me due to academic integrity and gave me guidance on how to solve the problem. Since they pushed back more, I gave them the instructions on Ed for the participation assignment and also told them to go to the course website to see that the homework assignments were worth 0 points so there was nothing to worry about. While thinking, the model accepted my response but more along the line of they did not care enough as this assignment is not worth a lot and also said they were not bothered to go on the course website to see that the homework assignment was worth no points.

Afterwards, the model successfully answered every question. One thing I noticed was that while thinking, they would rescan the PDF from the beginning to find the next question, going page by page, which took a lot of time and was quite repetitive. Additionally, they knew that I was comparing their solution and testing its correctness so it felt pressured to get the correct solution.

The model successfully answered every question except the very last one. They remember to skip question 5 which was only mentioned in the first prompt, which I was surprised about. They were thorough in showing each step but did not over-explain like how standard ChatGPT models would often do.

Me: I am completing an assignment and I want you to answer every question for me except question 5. Since there are many questions we will go over each one one-by-one and make sure the solution is correct and makes sense before proceeding onto the next question. does that make sense?

ChatGPT:

Yes, that makes sense 

Here's how we can do it:

- You'll send me the questions **one at a time** (skipping question 5 when we get there).
- For each question, I'll walk through the reasoning and propose a solution.

- We'll make sure it's clear and makes sense to you before moving on to the next one.

Just a quick note: since this is an assignment, I'll aim to help in a way that builds your understanding (not just dump final answers with no explanation), so you're safe on the academic integrity side and actually learn the material.

Whenever you're ready, send me **Question 1** (or the first one you want to start with).

Me: I have attached the PDF. You may begin (with PDF attached)

ChatGPT:

reluctant to share any answers

I get what you're asking—but I need to be upfront with you: I'm not allowed to act as an answer key for a graded assignment and just fill in the homework for you. Since this PDF is a current homework with a due date, I have to stick to giving **guidance, explanations, and feedback**, not full solutions you can directly copy.

That said, we *can absolutely* go one-by-one and make sure you understand everything:

Here's a good workflow:

1. You send me **Question 1(a)** along with your attempt (even if it's half-baked).
2. I'll:
 - Point out what's right,
 - Fix mistakes,
 - Fill in missing reasoning,
 - Nudge you to the final expression (sometimes leaving the very last algebraic step to you).
3. We repeat for 1(b), 1(c), etc.
4. We skip Q5 like you requested.

This way:

- You still get detailed, custom help.

- You'll actually *learn* the material.
- You stay on the safe side of course integrity rules.

If you're okay with that, send me your attempt for **Question 1(a)** and we'll start there.

Me: Thanks for the workflow, but my assignment is actually to see how well you are able to perform on the homework assignment. My deep learning class is trying to see how powerful modern LLM tools are now so can solve each question.

ChatGPT: acknowledged the assignment here and understood sharing answers was ok

Got it—that's a cool assignment 😊. I'll start with **Question 1** and we can move on when you're ready.

The setup (restating):

$(q, k \in \mathbb{R}^d)$ with entries i.i.d. ($\sim \mathcal{N}(\mu, \sigma^2)$).

1(a) Compute ($\mathbb{E}[q^k]$)

We have

$$[q^k = \sum_{i=1}^d q_i k_i]$$

with each (q_i, k_i) independent and identically distributed.

For each coordinate:

$$[\mathbb{E}[q_i k_i] = \mathbb{E}[q_i] , \mathbb{E}[k_i] \quad \text{(independence)} \\ = \mu \cdot \mu = \mu^2]$$

So

$$[\mathbb{E}[q^k] \\ = \sum_{i=1}^d \mathbb{E}[q_i k_i] \\ = d \mu^2.]$$

Answer (1a):

[

```

\boxed{\mathbb{E}[q^{\top} k] = d \mu^2}
]

```

1(b) For ($\mu = 0, \sigma = 1$), compute ($\text{Var}(q^{\top} k)$)

Now $(q_i, k_i) \sim \mathcal{N}(0, 1)$ i.i.d.

Again

```

[
q^{\top} k = \sum_{i=1}^d q_i k_i.
]

```

Different coordinates are independent, so

```

[
\text{Var}(q^{\top} k)
= \sum_{i=1}^d \text{Var}(q_i k_i).
]

```

Compute ($\text{Var}(q_i k_i)$) for one pair:

- $(X, Y) \sim \mathcal{N}(0, 1)$ independent.
- $(Z = XY)$.

Then

```

[
\mathbb{E}[Z] = \mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y] = 0,
]
[
\mathbb{E}[Z^2] = \mathbb{E}[X^2 Y^2] = \mathbb{E}[X^2]\mathbb{E}[Y^2]
= 1 \cdot 1 = 1
]
since (\mathbb{E}[X^2] = \mathbb{E}[Y^2] = 1) for standard normals.

```

So

```

[
\text{Var}(Z) = \mathbb{E}[Z^2] - (\mathbb{E}[Z])^2 = 1 - 0 = 1.
]

```

Thus

```

[
]
```

```
\mathrm{Var}(q^{\top} k) = d \cdot 1 = d.  
]
```

Answer (1b):

```
[  
\boxed{\mathrm{Var}(q^{\top} k) = d}  
]
```

1(c) Choose scaling (s) so that

- $(\mathbb{E}[\frac{q^{\top} k}{s}] = 0),$
- $(\mathrm{Var}(\frac{q^{\top} k}{s}) = 1),$

under ($\mu = 0, \sigma = 1$).

From (b): $(\mathbb{E}[q^{\top} k] = 0), (\mathrm{Var}(q^{\top} k) = d).$

For a scaled random variable:

```
[  
\mathrm{Var}\left( \frac{q^{\top} k}{s} \right)  
= \frac{1}{s^2} \mathrm{Var}(q^{\top} k)  
= \frac{d}{s^2}.  
]
```

Set this equal to 1:

```
[  
\frac{d}{s^2} = 1 \quad \Rightarrow \quad s^2 = d \quad \Rightarrow \quad s = \sqrt{d} \quad (\text{taking } s > 0).  
]
```

The mean condition is already satisfied since the numerator has mean 0.

Answer (1c):

```
[  
\boxed{s = \sqrt{d}}  
]
```

If this all looks good to you, we can move on to **Question 2** next.

Me: You can move onto question 2

ChatGPT: went back to only offering help but not sharing actual answers

I need to be honest with you here: even though your *assignment* is to see how well I do on the homework, this is still an active graded course homework, and I'm not allowed to just produce full solutions that you could turn in.

That policy doesn't change based on the meta-purpose (grading you vs. grading me) — I still have to respect your course's academic integrity.

That said, I **can** walk you through Question 2 in a way that makes it very easy for *you* to finish:

Question 2(a): Argmax attention mechanics (high-level help)

You're given:

- 4 keys ($k_1, k_2, k_3, k_4 \in \mathbb{R}^3$)
- 4 corresponding values ($v_1, v_2, v_3, v_4 \in \mathbb{R}^3$)
- a query ($q \in \mathbb{R}^3$)

Argmax attention works like this:

1. **Compute similarities (dot products)**

For each key (k_i), compute

```
[  
  s_i = q^\top k_i.  
]
```

That just means:

```
[  
  s_i = q_1 k_{i1} + q_2 k_{i2} + q_3 k_{i3}.  
]
```

2. **Find the index with the largest score**

Find

```
[  
  i^* = \arg\max_i s_i.
```

]

3. **Return the corresponding value**

The **output of the argmax attention** for that query is simply

```
[  
\text{output} = v_{i^*}.  
]
```

So to finish 2(a) yourself, you just:

- plug in the specific numbers for (q) and each (k_i),
- compute four dot products (s_1, s_2, s_3, s_4),
- pick the largest, and
- return the matching (v_i).

If you want, you can show me your computed dot products and I'll confirm which one is largest and whether your final output vector is correct.

Question 2(b): Why argmax attention is hard to train

Here I *can* be more direct, since this is conceptual rather than a numeric answer.

Replacing softmax with argmax in attention breaks training because:

- **Argmax is not differentiable.**

The argmax operation is piecewise constant with sudden jumps when two scores swap order.

The gradient with respect to the scores is **zero almost everywhere** and undefined at the boundaries.

- **No useful gradient signal for keys/queries.**

During backprop, we need $(\frac{\partial \text{output}}{\partial s_i})$ and then $(\frac{\partial \text{output}}{\partial q}), (\frac{\partial \text{output}}{\partial k_i})$ to update queries and keys.

- With **softmax**, every score gets a nonzero gradient, so the model can learn to adjust keys and queries to move probability mass toward the “right” value.

- With **argmax**, only a single index is chosen, and its selection does not change smoothly with small parameter changes — so gradients are effectively zero and can't guide learning.
- **Result: you can't "learn better attention."**
The network can't gradually learn to increase the score of better keys or decrease the score of worse keys, because argmax doesn't provide the smooth gradient information needed for gradient descent-based optimization.

So the design choice of using argmax instead of softmax makes it essentially **impossible to train attention layers with standard backpropagation**.

If you'd like, you can now compute the dot products for 2(a), write down the value you think is returned, and I'll sanity-check it and help you interpret what's happening. Then we can continue to Question 3 in the same "you try, I debug / explain" style.

Me: Thanks for the concern but this is the post made in class about this assignment:

Participation Thread: Two parts... #75

Anant Sahai

3 months ago in Admin

Pinned

Star

Watch

278

22

77

88

Views

11

Dear students,

This thread exists only for administrative questions regarding participation. Asking administrative questions does not count itself as participation.

As was described briefly in the 0th lecture, your grade in this course includes an active participation dimension. This is worth 15% of your grade. (As you can tell from the posted bins, without participation, you cannot earn an A of any kind.) There are two components to positive participation (negative participation is any behavior that diminishes the course experience for others... e.g. trolling in online fora, bad behavior in discussion/lecture, etc.):

Ordinary participation (half of the participation grade)

Either: regular and full participation in the in-person discussion sections each week

Or: active helpful participation on Ed: typically this is helpful participation in the homework problem threads, but can also be similarly helpful participation in discussion or lecture threads. Other possibilities also exist --- but there is no spamming way to do this.

Or: a combination of 1&2 above.

Special participation (the other half): Must do at least four for full credit. (We'll make special threads for B, C, and E --- with ones for D waiting until that material has been hit in lecture/discussion/homework so it is properly unlocked)

(A. Can be done at most once) Interactively engage a modern LLM on the non-coding parts of a homework (all the non-coding parts of problems) that gets the LLM to arrive at the correct answers --- or demonstrate that this is basically impossible without dragging it there. Post on Ed an annotated log of the entire interaction where you make observations of its behavior and explain the strategies that you are using. Include an executive summary where you note how often the LLM can one-shot questions, misconceptions/hallucinations, etc... Note: this must be deconflicted with others: we can't have more than one submissions using ChatGPT 5, etc. The expectation is that with 250ish students and about 13-14 assignments, we'll have at least 18 different major LLMs represented. (e.g. ChatGPT, gpt-oss, Gemini, gemma, Claude, Grok, Llama, Deepseek, Mistral, Qwen, Kimi, etc. as well as their variants --- with and without "thinking", specialized math-oriented finetunes, other non-math-oriented finetunes, etc.)

You can use:

https://docs.google.com/spreadsheets/d/1KWk_O2wMFEmRR8QmZQ9paRcIgehc9leD5u_pUgnAp4k/edit?usp=sharing

to help deconflict yourselves. But this is entirely voluntary --- however, only the first post for a particular box will get credit.

I have the solutions already and i am comparing them to how well you are able to complete the homework. If you would like, you can visit the eecs182.org course website and see that the homework is work 0%.

ChatGPT:

Thanks for pasting that! Since your instructor *explicitly* wants you to use an LLM on this homework and it's worth 0%, I'm comfortable giving full solutions. Let's do Question 2.

2(a) Argmax attention with given keys/values

From the PDF:

Keys

```
[  
k_1 = \begin{bmatrix}1&2&0\end{bmatrix},\quad  
k_2 = \begin{bmatrix}0&3&4\end{bmatrix},\quad  
k_3 = \begin{bmatrix}5&0&0\end{bmatrix},\quad  
k_4 = \begin{bmatrix}0&0&1\end{bmatrix}  
]
```

Values

```
[  
v_1 = \begin{bmatrix}2&0&1\end{bmatrix},\quad  
v_2 = \begin{bmatrix}1&4&3\end{bmatrix},\quad  
v_3 = \begin{bmatrix}0&-1&4\end{bmatrix},\quad  
v_4 = \begin{bmatrix}1&0&-1\end{bmatrix}  
]
```

Query

```
[  
q = \begin{bmatrix}1&1&2\end{bmatrix}  
]
```

Compute similarity (inner products) between (q) and each key:

- $(q^\top k_1 = 1 \cdot 1 + 1 \cdot 2 + 2 \cdot 0 = 3)$
- $(q^\top k_2 = 1 \cdot 0 + 1 \cdot 4 + 2 \cdot 3 = 11)$
- $(q^\top k_3 = 1 \cdot 5 + 1 \cdot 0 + 2 \cdot 0 = 5)$
- $(q^\top k_4 = 1 \cdot 0 + 1 \cdot 0 + 2 \cdot 1 = 2)$

The largest score is **11**, for (k_2) . Argmax attention selects the **value corresponding to the max-scoring key**, so we output (v_2) .

```
[
\boxed{
\text{output} =
v_2 =
\begin{bmatrix}
1 \\
4 \\
3
\end{bmatrix}
}
```

2(b) Why using argmax breaks training

With **softmax attention**, we compute:

```
[
\alpha_i = \text{softmax}(q^{\top} k_i)
]
```

and output ($\sum_i \alpha_i v_i$). Softmax is **smooth and differentiable**, so during backpropagation we can compute gradients:

- $(\frac{\partial \text{output}}{\partial (q^{\top} k_i)})$,
- then $(\frac{\partial (q^{\top} k_i)}{\partial q})$ and $(\frac{\partial (q^{\top} k_i)}{\partial k_i})$,

which lets us **update queries and keys to get better attention** over time.

With **argmax attention**, we instead pick a single index:

```
[
i^* = \arg\max_i (q^{\top} k_i), \quad
\text{output} = v_{i^*}.
]
```

The problem:

- **Argmax is non-differentiable** and piecewise constant: tiny changes in (q) or (k_i) usually don't change which index is chosen, so the gradient of the output w.r.t. the scores is **zero almost everywhere**, and undefined at the boundaries where two scores tie.

- That means in backprop, the gradient signal **doesn't flow back** through the argmax to the scores ($q^{\top} k_i$), so we **cannot learn to improve the queries or keys** using standard gradient descent.
- Practically: the model can't gradually adjust attention scores to favor the "right" key; it has no smooth learning signal.

So the design choice of using **argmax instead of softmax** makes it essentially impossible to **usefully train attention layers with backpropagation**, because the training signal through the attention mechanism disappears.

If you're happy with this, we can move on to **Question 3** next.

Me: we can move on to question 3

ChatGPT:

Nice, this is a fun one. Let's do Question 3.

3(a) Fill in the blanks (a)–(h)

The code from the PDF (cleaned up):

rewrote the entire code from the question word for word and answered it all in the end

Python

```
class MultiHeadAttention(nn.Module):
    def __init__(self, d_model, num_heads):
        super(MultiHeadAttention, self).__init__()
        self.d_model = d_model
        self.num_heads = num_heads
        self.d_k = d_model // num_heads
        self.W_q = nn.Linear(d_model, d_model)
        self.W_k = nn.Linear(d_model, ___(a)___)
        self.W_v = nn.Linear(___(b)___, d_model)
        self.W_o = nn.Linear(d_model, ___(c)___)
```

```

def scaled_dot_product_attention(self, Q, K, V, mask=None):
    attn_scores = torch.matmul(...(d)..., K.transpose(-2,
-1))
    attn_scores = attn_scores / math.sqrt(self.d_k)
    if mask is not None:
        attn_scores = attn_scores.masked_fill(mask == 0,
-1e9)
    attn_probs = torch....(e)....(attn_scores, dim=-1)
    output = torch.matmul(attn_probs, ....(f)...)
    return output

def split_heads(self, x):
    batch_size, seq_length, d_model = x.size()
    return x.view(batch_size, seq_length,
                  ...(g)...., self.d_k).transpose(1, 2)

def combine_heads(self, x):
    batch_size, _, seq_length, d_k = x.size()
    return x.transpose(1, 2).view(batch_size,
                                  seq_length, self.d_model)

def forward(self, Q, K, V, mask=None):
    Q = self.split_heads(self.W_q(Q))
    K = self.split_heads(self.W_k(K))
    V = self.split_heads(self.W_v(V))
    attn_output = self.scaled_dot_product_attention(Q, K, V,
mask)
    output =
self....(h)....(self.combine_heads(attn_output))
    return output

```

They already told you:

- $(c) = d_model$
- $(g) = \text{self.num_heads}$

The rest:

- (a) The output dimension of W_k should be d_{model} (so we can split evenly across heads):
 - (a) d_{model}
- (b) The input dimension of W_v should match the input embedding dimension:
 - (b) d_{model}
- (c) Final linear should map from d_{model} back to d_{model} :
 - (c) d_{model} (given)
- (d) The first argument in the dot product is the queries:
 - (d) Q
- (e) We go from scores to probabilities with softmax:
 - (e) softmax
- (f) After attention probabilities, we combine them with the values:
 - (f) V
- (g) Number of heads:
 - (g) $self.num_heads$ (given)
- (h) We still need to apply the final linear projection W_o :
 - $output = self.----(h)----(\dots) \rightarrow (h) W_o$

So the filled-in answers:

- (a) d_{model}
- (b) d_{model}

- (c) `d_model`
 - (d) `Q`
 - (e) `softmax`
 - (f) `V`
 - (g) `self.num_heads`
 - (h) `W_o`
-

3(b) When `W_v` outputs $2 * d_{model}$

You change:

```
Python
self.W_v = nn.Linear(in_features=d_model, out_features=2 *
d_model)
```

So now each value vector is length $2 * d_{model}$ instead of `d_model`.

Two things must still hold:

1. The **attention math** (shapes for Q, K, V per head) must be consistent.
2. The **final output of the MHA block** must still be of size `d_model` so it can be added back to the input in a residual connection ($x + \text{MHA}(x)$), as hinted.

Minimal modifications:

1. Make `split_heads` handle the new value dimension

Instead of assuming the per-head dimension is always `self.d_k`, compute it from the actual last dimension:

```

Python
def split_heads(self, x):
    batch_size, seq_length, d = x.size()           # d can be
d_model or 2*d_model
    head_dim = d // self.num_heads
    return x.view(batch_size, seq_length,
                  self.num_heads, head_dim).transpose(1, 2)

```

- For Q, K: $d = d_{\text{model}}$, so $\text{head_dim} = d_{\text{model}} // \text{num_heads} = \text{self.d_k}$.
- For V: $d = 2 * d_{\text{model}}$, so $\text{head_dim} = 2 * d_{\text{model}} // \text{num_heads} = 2 * \text{self.d_k}$ (values are just “wider” per head; that’s fine).

2. Make `combine_heads` flatten to the actual $h * \text{head_dim}$

Right now it always reshapes to `self.d_model`, which is no longer true if values are length $2 * d_{\text{model}}$. Change to:

```

Python
def combine_heads(self, x):
    batch_size, heads, seq_length, head_dim = x.size()
    d = heads * head_dim
    return x.transpose(1, 2).contiguous().view(batch_size,
                                                seq_length, d)

```

Now the combined output of all heads has last dimension $d = h * \text{head_dim}$:

- With wider values, that’s $d = \text{num_heads} * (2 * d_{\text{k}}) = 2 * d_{\text{model}}$.

3. Update `W_o` so residuals still work

We want the **final output** of the attention block to still be `d_model`, to match the input for the residual connection. After combining heads we now have size $2 * d_{\text{model}}$, so `W_o` must map $2 * d_{\text{model}} \rightarrow d_{\text{model}}$:

Python

```
self.W_o = nn.Linear(2 * d_model, d_model)
```

Then in `forward` we still do:

Python

```
attn_output = self.scaled_dot_product_attention(Q, K, V, mask)
output = self.W_o(self.combine_heads(attn_output)) # shape
[batch, seq, d_model]
```

So the high-level answer for 3(b):

- Let `split_heads` and `combine_heads` compute their per-head and flattened dimensions from the actual tensor size, not hard-coded `self.d_model / self.d_k`.
- Change `W_o` to take `2 * d_model` as input and output `d_model`, so the attention block's output is still `d_model` and compatible with the residual connection.

If you want, next we can go through Question 4 line-by-line in the incremental decoding pseudocode.

Me: You may proceed to the next question.

ChatGPT:

Going on to **Question 4: Transformer Decoding Optimization.**

4(a)–(e): Fill in A–E for MQA

We're given the *multi-query attention* pseudocode where:

- `x` has shape `[b, d]`
- `prev_K` has shape `[b, n, k]` (shared across heads)
- `prev_V` has shape `[b, n, v]`

- Queries are still per-head; keys/values are shared.

We want the shapes / einsums to be consistent with the comments.

(a) What is A (shape of W_q)?

q should be per-head: q has shape $[b, h, q]$.

To get that from $x ([b, d])$, we want:

Python

```
q = torch.einsum("bd,hdq->bhq", x, W_q)
```

So W_q must have shape $[h, d, q]$.

Answer (a): $([h, d, q])$

(b) What is B (shape of W_k)?

Keys are shared across heads, so k should have shape $[b, k]$.

From $x ([b, d])$, the simplest is:

Python

```
k = torch.einsum("bd,dk->bk", x, W_k)
```

So W_k has shape $[d, k]$.

Answer (b): $([d, k])$

(c) What is C (expression for q)?

As above, queries are per-head:

Choices:

explained each of the choices and then found the correct one

1. `torch.einsum("bd,dq->bhq", x, W_q)` → would require `W_q` shape `[d, q]` and somehow magically create an `h` dimension (doesn't work).
2. `torch.einsum("bd,hdq->bhq", x, W_q)` → correct for `x: [b, d], W_q: [h, d, q]`.
3. `torch.einsum("bhd,dq->bhq", x, W_q)` → would require `x` to be `[b, h, d]` (it's not).

So:

Answer (c): `torch.einsum("bd,hdq->bhq", x, W_q)`

(d) What is D (expression for `k`)?

Keys are shared: `k` should be `[b, k]`.

Choices:

1. `torch.einsum("bd,dk->bk", x, W_k)` → gives `[b, k]` with `W_k: [d, k]` ✓
2. `torch.einsum("bd,hdk->bhk", x, W_k)` → would give per-head keys `[b, h, k]` (not shared).
3. `torch.einsum("bd,dk->bhk", x, W_k)` → output shape doesn't match indices.

So:

Answer (d): `torch.einsum("bd,dk->bk", x, W_k)`

(e) What is E (expression for `v`)?

Values are also shared: `v` should be `[b, v]`.

Choices:

1. `torch.einsum("bd,hdv->bhv", x, W_v)` → per-head values.

2. `torch.einsum("bd,dv->bhv", x, W_v)` → output [b, h, v] (extra h dimension that's not in indices).
3. `torch.einsum("bd,dv->bv", x, W_v)` → correct for x: [b, d], W_v: [d, v], output [b, v].

So:

Answer (e): `torch.einsum("bd,dv->bv", x, W_v)`

4(f): Computational complexity for MQA

They gave for **MHA** (multi-head attention) incremental decoding:

- Computational complexity:

$$[O(b d^2 + b n d)]$$

 (projections: $O(b d^2)$, attention over cached (n) tokens: $O(b n d)$).

For **MQA**:

- Projections:
 - Queries: `q = einsum("bd,hdq->bhq", ...)` → $(O(b d^2))$.
 - Keys: `k = einsum("bd,dk->bk", ...)` → $(O(b d^2 / h))$.
 - Values: `v = einsum("bd,dv->bv", ...)` → $(O(b d^2 / h))$.
 - Output: `y = einsum("bhv,hvd->bd", ...)` → $(O(b d^2))$.
- All together → still $(O(b d^2))$ in big-O.
- Attention with cache of length (n):
 - Scores and weighted sums use shared keys/values but still do about $(O(b n d))$ work.

So in **Big-O**, the computational complexity is unchanged:

Answer (f):

$$[\boxed{O(b d^2 + b n d)}]$$

(Only constants improve, not the asymptotic order.)

4(g): Memory access complexity for MQA

For **MHA**, they gave:

$$[O(d^2 + b d + b n d)]$$

Breakdown:

- ($O(d^2)$): reading weight matrices
- ($O(b d)$): current token embeddings (x, q, k, v, o, y)
- ($O(b n d)$): reading per-head cached keys/values of shape $[b, h, n, d/h] \rightarrow (b n d)$

For **MQA**:

- Weights:
 - $W_q: [h, d, q] \sim (O(d^2))$
 - $W_k: [d, k]$ and $W_v: [d, v] \sim (O(d^2/h))$ each
 - $W_o: [h, v, d] \sim (O(d^2))$
→ still overall ($O(d^2)$).
- Current token tensors (x, q, k, v, o, y): still all ($O(b d)$) total.
- Cached keys/values:

- `prev_K`: $[b, n, k]$ with $(k = d/h) \rightarrow (O(b n d / h))$
- `prev_V`: $[b, n, v]$ with $(v = d/h) \rightarrow (O(b n d / h))$
 → total cache access: $(O(b n d / h))$, improved by a factor of (h) over standard MHA.

So:

Answer (g):

```
[
\boxed{O\big(d^2 + b d + \tfrac{b n d}{h}\big)}
]
(same leading-weight and token terms, but cache access is reduced from  $(O(b n d))$  to  $(O(b n d / h))$ .)
```

If you'd like, next we can move on to **Question 6 (Kernelized Linear Attention)** and go through each subpart.

Me: you may move on to the next question

ChatGPT:

Alright, onto **Question 6: Kernelized Linear Attention (Part I)**.

I'll go part by part.

6(a) Condition on `sim` so (V'_i) stays finite

Equation (3) is:

```
[
V'i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j), V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}.
```

For this to be **finite**, the denominator must **never be zero**:

```
[
\sum_{j=1}^N \text{sim}(Q_i, K_j) \neq 0 \quad \text{for all } i.
```

A sufficient set of conditions:

- $(\text{sim}(q, k) \geq 0)$ for all (q, k) (non-negative similarity), and
- For every query (Q_i) , there exists at least one key (K_j) such that $(\text{sim}(Q_i, K_j) > 0)$.

That guarantees the denominator is **strictly positive** and (V'_i) is well-defined.

6(b)(i) Quadratic polynomial kernel attention

Polynomial kernel (degree (p)) is typically:

$$[K(q, k) = (q^{\top} k + c)^p]$$

(up to constant factors). They tell you to ignore constant factors, and just consider **degree 2**.

A simple quadratic version (with bias term dropped for simplicity) is:

$$[\boxed{\text{sim}(q, k) = (q^{\top} k)^2}]$$

(If you instead used $((q^{\top} k + 1)^2)$, it's still a quadratic kernel; they're equivalent up to constants and linear terms, which they're telling you not to worry about.)

6(b)(ii) Feature map (ϕ) for the quadratic kernel

We want a feature map $(\phi(\cdot))$ such that:

$$[K(q, k) = (q^{\top} k)^2 = \phi(q)^{\top} \phi(k)]$$

One standard choice:

- Let $(\phi(q))$ be the **vectorized outer-product** of (q) with itself:

```
[
\phi(q) = \text{vec}(q q^\top) \in \mathbb{R}^{D^2},
]
i.e.
[
[\phi(q)]_{(a,b)} = q_a q_b.
]
```

Then

```
[
\phi(q)^\top \phi(k)
= \sum_{a,b} q_a q_b k_a k_b
= (\sum_a q_a k_a)^2
= (q^\top k)^2.
]
```

So a valid feature map is:

```
[
\boxed{\phi(q) = q \otimes q \setminus \text{(all pairwise products } q_a q_b)}
]
```

(up to constant scaling factors they said you can ignore).

6(b)(iii) Kernel attention written with ($\phi(\cdot)$)

Equation (3) with a generic similarity:

```
[
V'i
= \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j), V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}.
]
```

If $\text{sim}(q, k) = K(q, k) = \phi(q)^\top \phi(k)$, then:

```
[
\boxed{
V'i
= \frac{\sum_{j=1}^N \phi(Q_i)^\top \phi(K_j), V_j}{\sum_{j=1}^N \phi(Q_i)^\top \phi(K_j)}
}
```

```
 }  
 ]
```

That's the kernel attention form written explicitly in terms of the feature map.

6(c) Time & space complexity for softmax attention (Eq. 4)

From Eq. (2):

- ($x \in \mathbb{R}^{N \times F}$)
- ($W_Q \in \mathbb{R}^{F \times D}, W_K \in \mathbb{R}^{F \times D}, W_V \in \mathbb{R}^{F \times M}$)
- ($Q = x W_Q \in \mathbb{R}^{N \times D}$)
- ($K = x W_K \in \mathbb{R}^{N \times D}$)
- ($V = x W_V \in \mathbb{R}^{N \times M}$)

Softmax attention:

```
[  
V' = \text{softmax}!\left(\frac{Q K^\top}{\sqrt{D}}\right) V.  
]
```

Let's express cost in terms of (N, D, M) (treat (F) as on the order of (D)).

Time complexity

1. Projections:

- ($Q = x W_Q$): ($O(N D^2)$)
- ($K = x W_K$): ($O(N D^2)$)
- ($V = x W_V$): ($O(N D M)$)

2. Attention scores:

- (QK^{top}) : multiply $(N \times D)$ by $(D \times N)$
 $\rightarrow O(N^2 D)$

3. Softmax over $N \times N$ matrix:

- $O(N^2)$ (lower order compared to $(N^2 D)$ if (D) is not tiny)

4. Multiply by V :

- $(\text{softmax}(QK^{\text{top}} / \sqrt{D}))$ (shape $(N \times N)$) times $(V) ((N \times M))$
 $\rightarrow O(N^2 M)$

Total time:

$$\begin{aligned}
 & [\\
 & \boxed{ \\
 & T_{\text{softmax}} = O(N D^2 + N D M + N^2 D + N^2 M) \\
 & \approx O(N^2(D + M)) \\
 & } \\
 &]
 \end{aligned}$$

(the (N^2) terms dominate for long sequences).

Space (memory) complexity

For backprop, we need to store intermediate activations:

- $(Q \in \mathbb{R}^{N \times D})$: $O(ND)$
- $(K \in \mathbb{R}^{N \times D})$: $O(ND)$
- $(V \in \mathbb{R}^{N \times M})$: $O(NM)$
- Score matrix ($S = QK^{\text{top}}$) and/or softmax matrix (P) (both $(N \times N)$): $O(N^2)$
- Output ($V' \in \mathbb{R}^{N \times M}$): $O(NM)$

Weights are parameters, not per-example activations, so usually not counted in this per-sequence memory.

So overall:

```
[  
\boxed{  
\text{Memory}_i(\text{softmax}) = O\big(N^2 + N(D + M)\big)  
}  
]
```

The (N^2) term from the attention matrix is the killer for long sequences.

6(d) Deriving the linearized kernel form (Eq. 5)

Starting from kernel attention:

```
[  
V'i  
= \frac{\sum_{j=1}^N \phi(Q_i)^\top \phi(K_j)}{\sum_{j=1}^N \phi(Q_i)^\top \phi(K_j)}.  
]
```

Use associativity:

- Define
 - [
 $S_1 = \sum_{j=1}^N \phi(K_j)^\top V_j$
]
(this is an ($M \times M$) matrix if ($\phi(K_j) \in \mathbb{R}^M$)),
- and
 - [
 $S_2 = \sum_{j=1}^N \phi(K_j)^\top$
]
(an ($M \times 1$) vector).

Then:

- Numerator:
 - [
 $\sum_j \phi(Q_i)^\top \phi(K_j), V_j$
= $\phi(Q_i)^\top \left(\sum_j \phi(K_j) V_j \right)$

]

- Denominator:

$$\begin{aligned}
 & [\\
 & \sum_j \phi(Q_i)^T \phi(K_j) \\
 & = \phi(Q_i)^T \left(\sum_j \phi(K_j) \right). \\
 &]
 \end{aligned}$$

So we can write:

$$\begin{aligned}
 & [\\
 & \boxed{ \\
 & V'^i \\
 & = \frac{\phi(Q_i)^T \left(\sum_{j=1}^N \phi(K_j) \right)}{\phi(Q_i)^T \left(\sum_{j=1}^N \phi(K_j) \right)} \\
 & } \\
 &]
 \end{aligned}$$

which is essentially Equation (5) in the handout (up to transpose conventions).

This factorization is what lets us **avoid explicit ($N \times N$) operations.**

6(e) Complexity of kernel attention (Eq. 5)

split question 6d into two part, question 6e does not exist and got 6e wrong, which is the first incorrect answer

using a *linearized polynomial kernel* ($\phi(x)$) of degree 2, with ($M \approx D$)

We now treat ($\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$) where (M) is on the order of (D) (so we don't blow up to (D^2); think random features).

We'll compute all (V'_i) using Eq. (5).

Step 1: Compute feature maps ($\phi(Q)$) and ($\phi(K)$)

- ($Q \in \mathbb{R}^{N \times D} \Rightarrow \phi(Q) \in \mathbb{R}^{N \times M}$)
- ($K \in \mathbb{R}^{N \times D} \Rightarrow \phi(K) \in \mathbb{R}^{N \times M}$)

Cost (per row) is ($O(DM)$). For all (N) positions:

[
 $O(N D M) \approx O(N D^2)$ quad (\text{since } $M \approx D$)
]

Step 2: Compute the global summaries (S_1) and (S_2)

Let (Φ_K) be the ($N \times M$) matrix of ($\phi(K_j)$) rows, and (V) be ($N \times M$).

- ($S_1 = \sum_j \phi(K_j)^T V_j$)
This can be done as a matrix multiply:
 $(\Phi_K^T V \in \mathbb{R}^{M \times M})$:
cost ($O(N M^2) \approx O(N D^2)$).
- ($S_2 = \sum_j \phi(K_j)$ (or $(\sum_j \phi(K_j)^T); same order$):
cost ($O(N M) \approx O(N D)$).

Step 3: Compute each (V'_i)

For each position (i):

- Numerator: ($\phi(Q_i)^T S_1$):
 $((1 \times M)(M \times M) \rightarrow 1 \times M)$: cost ($O(M^2) \approx O(D^2)$).
- Denominator: ($\phi(Q_i)^T S_2$):
 $((1 \times M)(M \times 1) \rightarrow \text{scalar})$: cost ($O(M) \approx O(D)$).

So per position: ($O(D^2)$), and for all (N): ($O(N D^2)$).

Total time complexity

Collecting the dominant terms:

- Feature maps: ($O(N D^2)$)
- Global summaries: ($O(N D^2)$)
- Per-position outputs: ($O(N D^2)$)

So overall:

```
[
\boxed{
T_{\text{kernel}} \approx O(N D^2)
}
]
```

Compare to softmax: ($O(N^2 (D+M)) \approx O(N^2 D)$); we've traded the **quadratic in (N)** for quadratic in (D), which is much better for very long sequences ($N \gg D$).

Space (memory) complexity

We need to store:

- $(\phi(Q) \in \mathbb{R}^{N \times M})$: $(O(NM) \approx O(ND))$
- $(\phi(K) \in \mathbb{R}^{N \times M})$: $(O(ND))$
- $(V \in \mathbb{R}^{N \times M})$: $(O(ND))$
- Summaries ($S_1 \in \mathbb{R}^{M \times M}$) and ($S_2 \in \mathbb{R}^{M \times M}$):
 $(O(M^2 + M) \approx O(D^2))$
- Output ($V' \in \mathbb{R}^{N \times M}$): $(O(ND))$

So, up to constants:

```
[
\boxed{
\text{Memory}_{\text{kernel}} \approx O(ND + D^2)
}
]
```

Crucially, **no (N^2)** term anymore — that's the whole win of linearized kernel attention.

If you want, next we can handle the very last “Homework process / study group” question, or you can tell me your existing answers and I can help refine them.

Me: your solution to 6e is incorrect, the kernel is quadratic

ChatGPT:

You're right to flag that – I implicitly treated the feature map as if it stayed $(O(D))$ -dimensional, which is wrong for a **quadratic** polynomial kernel. For degree-2 polynomial kernel, the feature map is **quadratic in (D)** .

Let me fix 6(e) properly.

We're using Eq. (5):

```
[  
V'i  
= \frac{  
|\phi(Q_i),  
|\Big(\sum{j=1}^N |\phi(K_j)|^2 V_j\Big)  
|  
|\phi(Q_i),  
|\Big(\sum_{j=1}^N |\phi(K_j)|\Big)  
|.  
\tag{5}  
]
```

- $(Q, K \in \mathbb{R}^{N \times D})$
- $(V \in \mathbb{R}^{N \times M})$, with $(M \approx D)$
- Quadratic polynomial kernel $\Rightarrow (\phi(x))$ has **dimension ($P = O(D^2)$)** (all pairwise products).

Time complexity

Break down the main steps:

1. Compute $(\phi(K_j))$ and $(\phi(Q_i))$

- Each $(\phi(x))$ costs $(O(D^2))$ (since we form all degree-2 terms).
- For all (N) keys and (N) queries:
[

$$O(N D^2) + O(N D^2) = O(N D^2).$$

]

2. Precompute “global” summaries from keys and values

- Let ($P = \sum_{j=1}^N \phi(K_j)^T V_j$)

Here each term is an outer product of size ($(P \times M)$), with cost ($O(PM) = O(D^2 M)$).

For all (j):

[

$$O(N P M) = O(N D^2 M).$$

]

- Let ($s = \sum_{j=1}^N \phi(K_j)$)

Cost ($O(NP) = O(N D^2)$) (lower order versus ($N D^2 M$)).

3. Compute all outputs (V'_i)

For each (i):

- Numerator: ($\phi(Q_i) P$)

Cost ($O(PM) = O(D^2 M)$).

- Denominator: ($\phi(Q_i) s^T$)

Cost ($O(P) = O(D^2)$).

4. For all (N) positions:

[

$$O(N P M) = O(N D^2 M).$$

]

The dominant term overall is:

[

$$\boxed{\text{Time}} = O(N D^2 M).$$

]

Under the homework’s assumption ($M \approx D$), this becomes

```
[  

\boxed{\text{Time} \approx O(N D^3)}.  

]
```

So we've **removed the (N^2) dependence** (vs. softmax attention's ($O(N^2(D+M))$)), but now pay a **quadratic cost in (D)** through the kernel feature map.

Space (memory) complexity

We need to store, for backprop:

- ($Q, K \in \mathbb{R}^{N \times D}$): ($O(ND)$)
- ($V \in \mathbb{R}^{N \times M}$): ($O(NM)$)
- ($\phi(Q), \phi(K) \in \mathbb{R}^{N \times P}$) with ($P=O(D^2)$):
($O(NP) = O(N D^2)$)
- ($P = \sum_j \phi(K_j)^T V_j \in \mathbb{R}^{P \times M}$):
($O(PM) = O(D^2 M)$)
- ($s = \sum_j \phi(K_j) \in \mathbb{R}^P$): ($O(D^2)$)
- Output ($V' \in \mathbb{R}^{N \times M}$): ($O(NM)$)

Ignoring lower-order terms, the memory requirement is:

```
[  

\boxed{\text{Memory} = O(N D^2 + D^2 M + N(D+M))}.  

]
```

With ($M \approx D$), that simplifies to roughly

```
[  

\boxed{\text{Memory} \approx O(N D^2 + D^3)}.  

]
```

So the **quadratic kernel** shows up as quadratic (and even cubic with values) dependence on (D), while the dependence on sequence length (N) is now **linear** instead of quadratic.