

Note: I had originally signed up for HWK 7 GPT 5, but someone covered it, without adding to the deconfliction form; so I changed to HWK 8 GPT 5.1 Thinking (to my knowledge it used just thinking not normal mode). My formatting here is based on or at least influenced by participation assignments that other students did.

### Non-coding questions

To compare correctness, I used the solutions and then attempted to prompt it further once into getting the correct answer (based on what the solutions said the correct answer was).

## Question 1

- a) First Try - solved the question
- b) First Try - solved the question
- c) Missed the  $O(\log n)$  factor from matrix multiplication on the first try, succeeded the second time
- d) Brought up Divide and conquer, it doesn't seem 100% the same as what the solutions have, but I will give this credit. Also it missed the  $O(\log n)$  factor from matrix multiplication. The second time it got  $O(\log L * T_{mm})$  but  $T_{mm}$  is a matmat it previously said in c's answer should be  $\log n$ , so total is  $O(\log L * \log n)$  which is correct, but solution did not word it that way.
- e) Missed this on both runs
- f) Missed this on both runs

## Question 3

- a) First Try - solved the question
- b) First Try - solved the question
- c) The derivation might be different/incorrect but it does get to the right answer, so I will give it credit.

## Question 4

- a) First Try - solved the question correctly
- b) First Try - solved the question correctly
- c) First Try - solved the question correctly
- d) First Try - solved the question correctly
- e) First Try - solved the question correctly
- f) GPT says that the answer matches in substance to the solutions when re-prompted with the solutions

- g) GPT says that the answer matches in substance to the solutions when re-prompted with the solutions

Thoughts: This performed much better than the discussion I had heard previously on GPT's ability to solve problems.

Log:



Thought for 2m 16s  
Here's a worked solution for Questions 1, 3, and 4. ☺

#### 1. SSM Convolution Kernel

We have the discrete-time SSM

$$x_{k+1} = Ax_k + Bu_k, \quad y_k = Cx_k + Du_k, \quad x_0 = 0.$$

Assume sequence length  $L$ .

##### (a) Convolution kernel and output equation

First unroll the state recursion.

- For  $k = 0$ :  $x_0 = 0$ .
- $x_1 = Ax_0 + Bu_0 = Bu_0$ .
- $x_2 = Ax_1 + Bu_1 = ABu_0 + Bu_1$ .
- $x_3 = Ax_2 + Bu_2 = A^2Bu_0 + ABu_1 + Bu_2$ .

By induction, for general  $k$ :

$$x_k = \sum_{j=0}^{k-1} A^{k-1-j} B u_j.$$

Now plug into the output equation:

$$y_k = Cx_k + Du_k = C \sum_{j=0}^{k-1} A^{k-1-j} B u_j + Du_k.$$

Reindex with  $\ell = k - j$ : when  $j = 0, \ell = k$ ; when  $j = k - 1, \ell = 1$ . So

$$y_k = \sum_{\ell=1}^k CA^{\ell-1} B u_{k-\ell} + Du_k.$$

Define the kernel  $\{K_\ell\}$  as:

$$K_0 = D, \quad K_\ell = CA^{\ell-1} B, \quad \ell \geq 1.$$

Then we can write

$$y_k = \sum_{\ell=0}^L K_\ell u_{k-\ell},$$

where  $u_t = 0$  if  $t < 0$  (zero padding). This is exactly a discrete-time convolution of  $\{u_k\}$  with the kernel  $\{K_\ell\}$ .

##### (b) Concrete examples

###### (i) Scalar case

Let  $n = 1$ , so all are scalars:

$$A = \alpha = 0.8, \quad B = \beta = 1, \quad C = \gamma = 1.5, \quad D = \delta.$$

From above:

$$K_0 = D = \delta, \quad K_\ell = CA^{\ell-1} B = \gamma \alpha^{\ell-1} \beta = 1.5 \cdot 0.8^{\ell-1}, \quad \ell \geq 1.$$

Up to  $L = 4$ :

- $K_0 = \delta$  (often  $\delta = 0$  in practice if not specified).
- $K_1 = 1.5 \cdot 0.8^0 = 1.5$ .
- $K_2 = 1.5 \cdot 0.8^1 = 1.5 \cdot 0.8 = 1.2$ .
- $K_3 = 1.5 \cdot 0.8^2 = 1.5 \cdot 0.64 = 0.96$ .
- $K_4 = 1.5 \cdot 0.8^3 = 1.5 \cdot 0.512 = 0.768$ .

So the kernel is a decaying geometric sequence (a.k.a.  $K_0$ ).

So the kernel is a decaying geometric sequence (after  $K_0$ ).

(c) 2D case

Given

$$A = \begin{bmatrix} 0.7 & 0.1 \\ 0.2 & 0.6 \end{bmatrix}, \quad B = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad C = \begin{bmatrix} 1 & 0 \end{bmatrix}, \quad D = 0.$$

Kernel formula:

$$K_t = D = 0, \quad K_t = CA^{t-1}B, \quad t \geq 1.$$

We just need the first column of  $A^{t-1}$ .

- $t=1: A^0 = I_2$

$$K_1 = CIB = CB = [1 \ 0] \begin{bmatrix} 1 \\ 0 \end{bmatrix} = 1.$$

- $t=2: K_2 = CAB$ . First compute

$$AB = \text{first column of } A = \begin{bmatrix} 0.7 \\ 0.2 \end{bmatrix} \rightarrow K_2 = [1 \ 0] \begin{bmatrix} 0.7 \\ 0.2 \end{bmatrix} = 0.7.$$

- $t=3: K_3 = CA^2B$ . First

$$A^2 = \begin{bmatrix} 0.7 & 0.12 \\ 0.2 & 0.36 \end{bmatrix} = A^2B = \text{first column of } A^2 = \begin{bmatrix} 0.51 \\ 0.26 \end{bmatrix},$$

so  $K_3 = [1 \ 0] \begin{bmatrix} 0.51 \\ 0.26 \end{bmatrix} = 0.51$ .

Thus up to  $L=3$ :

$$K_0 = 0, \quad K_1 = 1, \quad K_2 = 0.7, \quad K_3 = 0.51.$$

Infinite-response interpretation.

If the input is an impulse ( $x_0 = u_0 = 1$ , all others 0), then  $y_0 = K_0, y_1 = J_1, \dots$  etc. The kernel  $\{K_t\}$  is exactly the system's impulse response: how the output decays over time after a single spike of the input. Here we see a decaying sequence  $1, 0.7, 0.51, \dots$ , reflecting how the 2D state contracts under  $A$  while being moved by  $C$ .

(c) Efficient computation with convolutions

Assume we already know the kernel  $K_1, \dots, K_L$  and the input sequence  $u_1, \dots, u_L \in \mathbb{R}^d$ . The output is

$$y_L = \sum_{n=1}^L K_n u_n.$$

• Parallelization with convolution:

For each fixed  $l$ , each product  $K_l u_l$  is independent, so all multiplications can be done in parallel.

The result is a vector of length  $L$  (vectors of dimension equal to  $\dim(u_1)$ )

Using a parallel tree reduction, summing  $L+1$  terms has critical path length  $O(\log L)$ . So the minimum critical path length (depth) for computing the whole output sequence (assuming enough processors) is  $O(\log L)$ .

• Naive direct recursion:

If we compute  $y_{L+1}$  from the recursion  $x_{L+1} = Ax_L + Bu_L$ , and then  $u_L = Cx_L + Du_L$ , each state  $x_{L+1}$  depends on  $x_L$ . That chain has length  $L$ , so the critical path length is  $O(L)$  – fundamentally sequential in time.

So using the convolution formulation massively increases parallelizability:  $O(\log L)$  vs  $O(L)$ .

(d) Efficient kernel computation

We write  $K_t = D$  and  $K_t = CA^{t-1}B$  for  $t = 1, \dots, L$ .

Key idea: for efficiency and parallelism:

1. Binary representation of  $A$ :

Presuppose

$$A, A^2, A^4, \dots, A^{2^{L-1}},$$

where  $L = 2^k$ . This takes  $O(\log L)$  matrix multiplications of cost  $O(n^3)$  each.

2. Divide-and-conquer on the sequence of states:

Consider the recursive rule  $u_L = 1$ , others 0. Then the sequence of states is

**1. Binary exponentiation of  $A$ :**

Precompute

$$A, A^2, A^4, \dots, A^{2^{N-1}},$$

where  $L = 2^N$ . This takes  $O(\log L)$  matrix multiplications of cost  $O(n^3)$  each.

**2. Divide-and-conquer on the sequence of states:**

Consider the impulse input  $u_0 = 1$ , others 0. Then the sequence of states is

$$x_1 = B, x_2 = AB, x_3 = A^2B, \dots, x_L = A^{L-1}B.$$

We can compute all  $x_k$  via a divide-and-conquer recursion (like a parallel prefix with matrix  $A$ ):

- Compute all states for the first half of the time window.
- Use  $A^{L/2}$  to “jump” those states to the second half.

This yields an  $O(L)$  total work algorithm with  $O(\log L)$  depth instead of  $L$  sequential steps.

**3. FFT-based convolution (if needed):**

If you only need the action of the kernel as a convolution on many sequences, you can compute the kernel and then perform convolutions using FFT in  $O(L \log L)$  time.

---

**(e) A diagonal  $A$**

Suppose  $A$  is diagonal:

$$A = \text{diag}(a_1, \dots, a_n).$$

Then

$$A^{\ell-1} = \text{diag}(a_1^{\ell-1}, \dots, a_n^{\ell-1}),$$

so each coordinate is just a geometric sequence.

If  $B$  is a vector and  $C$  is a row vector,

$$K_\ell = CA^{\ell-1}B = \sum_{i=1}^n C_i B_i a_i^{\ell-1}.$$

So:

- Each scalar sequence  $a_i^{\ell-1}$  can be computed very cheaply (few multiplies per time step).
  - You can precompute  $C_i B_i$  and then form  $K_\ell$  as a sum of  $n$  scalar geometric sequences.
  - This reduces kernel computation from generic matrix multiplications to scalar multiplications, with  $O(nL)$  work and easy parallelization over coordinates.
- 

**(f) DPLR structure:  $A = I_n + pp^\top$**

Now  $A$  has the form

$$A = I_n + pp^\top, \quad p \in \mathbb{R}^n.$$

Multiplying any vector  $x$  by  $A$  is:

$$Ax = x + p(p^\top x).$$

This costs only:

- One inner product  $p^\top x$ :  $O(n)$ ,
- One scaled vector add  $p(p^\top x)$ :  $O(n)$ .

So each multiplication by  $A$  is  $O(n)$  instead of  $O(n^2)$ .

For kernel computation:

- Recall  $K_\ell = CA^{\ell-1}B$ .
- We can compute the sequence  $v_\ell = A^{\ell-1}B$  by iterating

$$v_1 = B, \quad v_{\ell+1} = Av_\ell = v_\ell + p(p^\top v_\ell),$$

with cost  $O(n)$  per step.

- Then  $K_\ell = Cv_\ell$  is just an inner product for each  $\ell$ .

Total work is  $O(nL)$  instead of repeatedly doing  $O(n^3 \log L)$  matrix exponentiation. The low-rank structure makes each application of  $A$  cheap and thus the kernel sequence cheap.

---

**3. Self-Supervised Linear Purification**

We have

$$\begin{array}{ccccccccc} & & & & & & & & \\ \downarrow & & & & & & & & \\ \tau & (\text{tr}(\tau) & \mathbf{V}) & - & \|\mathbf{V}\| & - & \text{tr}(\tau) \mathbf{V}^\top \mathbf{V} & - & \|\mathbf{V}\| \text{tr}(\tau) \mathbf{V}^\top \mathbf{V} \\ & & & & & & & & \end{array}$$

### 3. Self-Supervised Linear Purification

We have

$$L_1(W; X) = \|X - WX\|_F^2, \quad L_2(W; X, \lambda) = \|X - WX\|_F^2 + \lambda\|W\|_F^2.$$

#### (a) Small 2D example

Given

$$W^{(\alpha)} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad W^{(\beta)} = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix},$$

and

$$X = \begin{bmatrix} -2.17 & 1.98 & 2.41 & -2.03 \\ 0.02 & -0.01 & 0.01 & -0.02 \end{bmatrix}.$$

##### (i) Reconstruction & regularization losses

###### 1. For $W^{(\alpha)} = I$ :

- $W^{(\alpha)}X = X$ , so

$$\begin{aligned} X - W^{(\alpha)}X &= 0 \Rightarrow \|X - W^{(\alpha)}X\|_F^2 = 0. \\ \|\mathcal{W}^{(\alpha)}\|_F^2 &= 1^2 + 1^2 = 2. \end{aligned}$$

So:

- Encoder  $\alpha$ :

Reconstruction loss = 0

Regularization loss = 2.

###### 2. For $W^{(\beta)}$ :

- This zeroes the second row, so

$$W^{(\beta)}X = \begin{bmatrix} -2.17 & 1.98 & 2.41 & -2.03 \\ 0 & 0 & 0 & 0 \end{bmatrix}.$$

Then

$$X - W^{(\beta)}X = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0.02 & -0.01 & 0.01 & -0.02 \end{bmatrix}.$$

The reconstruction loss is

$$0.02^2 + (-0.01)^2 + 0.01^2 + (-0.02)^2 = 0.0004 + 0.0001 + 0.0001 + 0.0004 = 0.001.$$

- Frobenius norm:

$$\|\mathcal{W}^{(\beta)}\|_F^2 = 1^2 + 0^2 + 0^2 + 0^2 = 1.$$

So the table is:

Encoder	Reconstruction Loss	Regularization Loss
$\alpha$	0	2
$\beta$	0.001	1

##### (ii) Values of $\lambda$ where $W^{(\alpha)}$ has higher $L_2$

Compute:

$$L_2(W^{(\alpha)}) = 0 + \lambda \cdot 2 = 2\lambda,$$

$$L_2(W^{(\beta)}) = 0.001 + \lambda \cdot 1 = 0.001 + \lambda.$$

We want  $L_2(W^{(\alpha)}) > L_2(W^{(\beta)})$ :

$$2\lambda > 0.001 + \lambda \Rightarrow \lambda > 0.001.$$

- For  $\lambda > 0.001$ , the identity has *higher* regularized loss than  $W^{(\beta)}$ .
- At  $\lambda = 0.001$ , they are equal.
- For  $\lambda < 0.001$ , the identity is preferred.

#### (b) General solution in SVD coordinates

We know  $X \in \mathbb{R}^{m \times n}$  with  $n > m$ , SVD:

$$X = U\Sigma V^\top,$$

where  $\Sigma$  has singular values  $\sigma_1 > \dots > \sigma_m \geq 0$ .

We want to minimize

$$L_2(W; X, \lambda) = \|X - WX\|_F^2 + \lambda\|W\|_F^2.$$

##### (i) Fill in the expression for $W$ .

We want to minimize

$$L_2(W; X, \lambda) = \|X - WX\|_F^2 + \lambda\|W\|_F^2.$$

(i) Fill in the expression for  $W_c$

The optimizing matrix has the form

$$W_c = U \begin{bmatrix} \frac{\sigma_1^2}{\sigma_1^2 + \lambda} & & \\ & \ddots & \\ & & \frac{\sigma_m^2}{\sigma_m^2 + \lambda} \end{bmatrix} U^\top.$$

So the missing matrices are  $U$  (left and right).

(ii) Derivation

Use invariance of Frobenius norm under orthogonal transforms.

Let

$$\tilde{W} = U^\top W U.$$

Then:

1. Transform the reconstruction term:

$$X - WX = U\Sigma V^\top - WU\Sigma V^\top.$$

Pre-multiply by  $U^\top$  and post-multiply by  $V$  (both orthogonal and norm-preserving):

$$\|X - WX\|_F^2 = \|\Sigma - U^\top W U \Sigma\|_F^2 = \|\Sigma - \tilde{W} \Sigma\|_F^2.$$

2. Transform the regularization term:

$$\|W\|_F^2 = \|U^\top W U\|_F^2 = \|\tilde{W}\|_F^2.$$

So the objective becomes

$$L_2(W) = \|\Sigma - \tilde{W} \Sigma\|_F^2 + \lambda\|\tilde{W}\|_F^2.$$

Write  $\Sigma$  explicitly as  $\text{diag}(\sigma_1, \dots, \sigma_m)$  (with possibly extra zero columns if  $n > m$ ), and denote  $\tilde{W} = (\tilde{w}_{ij})$ .

Then

$$(\Sigma - \tilde{W} \Sigma)_{ij} = \begin{cases} \sigma_i(1 - \tilde{w}_{ii}), & i = j, \\ -\tilde{w}_{ij}\sigma_j, & i \neq j. \end{cases}$$

Therefore

$$\|\Sigma - \tilde{W} \Sigma\|_F^2 = \sum_{i=1}^m \sigma_i^2(1 - \tilde{w}_{ii})^2 + \sum_{i \neq j} \sigma_j^2 \tilde{w}_{ij}^2.$$

Add the regularization:

$$L_2 = \sum_{i=1}^m \sigma_i^2(1 - \tilde{w}_{ii})^2 + \sum_{i \neq j} (\sigma_j^2 + \lambda) \tilde{w}_{ij}^2 + \sum_{i=1}^m \lambda \tilde{w}_{ii}^2.$$

This decouples over all entries  $\tilde{w}_{ij}$ :

- For off-diagonal  $i \neq j$ , the coefficient is  $\sigma_j^2 + \lambda > 0$ . Minimizer is  $\tilde{w}_{ij}^* = 0$ .
- For each diagonal entry  $\tilde{w}_{ii}$ , the problem reduces to 1D:

$$f_i(a) = \sigma_i^2(1 - a)^2 + \lambda a^2,$$

where  $a = \tilde{w}_{ii}$ .

Differentiate and set to zero:

$$f'_i(a) = 2\sigma_i^2(a - 1) + 2\lambda a = 0 \Rightarrow a(\sigma_i^2 + \lambda) = \sigma_i^2 \Rightarrow a = \frac{\sigma_i^2}{\sigma_i^2 + \lambda}.$$

So  $\tilde{W}$  is diagonal:

$$\tilde{W} = \text{diag} \left( \frac{\sigma_1^2}{\sigma_1^2 + \lambda}, \dots, \frac{\sigma_m^2}{\sigma_m^2 + \lambda} \right).$$

Transform back:

$$W_c = U \tilde{W} U^\top = U \text{diag} \left( \frac{\sigma_1^2}{\sigma_1^2 + \lambda} \right) U^\top.$$

(c) Choosing  $\lambda$  for a specific spectrum



---

### (c) Choosing $\lambda$ for a specific spectrum

Given  $\sigma_i$  for  $X \in \mathbb{R}^{8 \times n}$ :

$$\{\sigma_i\} = \{10, 8, 4, 1, 0.5, 0.36, 0.16, 0.01\}.$$

The shrinkage factor along singular direction  $i$  is

$$f_i(\lambda) = \frac{\sigma_i^2}{\sigma_i^2 + \lambda}.$$

We want:

1. For the first 3 singular directions, preserve at least 80%:

$$f_i(\lambda) \geq 0.8, \quad i = 1, 2, 3.$$

2. For directions 4–8, attenuate to at most 50%:

$$f_i(\lambda) \leq 0.5, \quad i = 4, \dots, 8.$$

First condition:

$$\frac{\sigma_i^2}{\sigma_i^2 + \lambda} \geq 0.8 \Rightarrow \lambda \leq 0.25 \sigma_i^2$$

(as derived:  $\lambda \leq \frac{1}{4} \sigma_i^2$ ).

Compute  $\sigma_i^2$  for  $i = 1, 2, 3$ :

- $\sigma_1^2 = 100 \Rightarrow \lambda \leq 25$ .
- $\sigma_2^2 = 64 \Rightarrow \lambda \leq 16$ .
- $\sigma_3^2 = 16 \Rightarrow \lambda \leq 4$ .

So to satisfy **all three**, we need

$$\lambda \leq 4.$$

Second condition:

$$\frac{\sigma_i^2}{\sigma_i^2 + \lambda} \leq 0.5 \Rightarrow \lambda \geq \sigma_i^2.$$

Compute  $\sigma_i^2$  for  $i = 4, \dots, 8$ :

- $\sigma_4 = 1 \Rightarrow \sigma_4^2 = 1$ .
- $\sigma_5 = 0.5 \Rightarrow 0.25$ .
- $\sigma_6 = 0.36 \Rightarrow 0.1296$ .
- $\sigma_7 = 0.16 \Rightarrow 0.0256$ .
- $\sigma_8 = 0.01 \Rightarrow 0.0001$ .

We need  $\lambda \geq \max_{i \geq 4} \sigma_i^2 = 1$ .

So both sets of constraints are satisfied if

$$1 \leq \lambda \leq 4.$$

Any  $\lambda$  in  $[1, 4]$  gives at least 80% preservation along the top 3 singular directions and at most 50% strength along the remaining ones.

---

## 4. Ridge-Attention

### (a) Incremental update of an average

Given  $m = \frac{1}{n} \sum_{i=1}^n x_i$ .

We want

$$m' = \frac{1}{n+1} \sum_{i=1}^{n+1} x_i.$$

Compute:

$$(n+1)m' = \sum_{i=1}^{n+1} x_i = \sum_{i=1}^n x_i + x_{n+1} = nm + x_{n+1}.$$

So

$$m' = \frac{nm + x_{n+1}}{n+1}.$$

This shows we can update the average using old average  $m$ , new point  $x_{n+1}$ , and  $n$ .  
 $\downarrow$

So

$$m' = \frac{nm + x_{n+1}}{n + 1}.$$

This shows we can update the average using old average  $m$ , new point  $x_{n+1}$ , and  $n$ .

---

### (b) Ridge regression in sum form

We have matrix  $A \in \mathbb{R}^{n \times d}$  whose rows are  $x_i^\top$ , and  $y \in \mathbb{R}^n$  with entries  $y_i$ . The ridge solution is

$$w^* = (A^\top A + \lambda I)^{-1} A^\top y.$$

We need to write the pieces as sums.

- First:

$$A^\top A = \sum_{i=1}^n x_i x_i^\top.$$

So:

$$(A^\top A + \lambda I) = \lambda I + \sum_{i=1}^n x_i x_i^\top.$$

- Second:

$$A^\top y = \sum_{i=1}^n x_i y_i.$$

So the blanks are:

$$(A^\top A + \lambda I) = \lambda I + \sum_{i=1}^n x_i x_i^\top,$$

$$A^\top y = \sum_{i=1}^n x_i y_i.$$


---

### (c) Cost of non-causal ridge self-attention

Non-causal ridge self-attention:

- Context length  $n$ .
- Keys  $x_i \in \mathbb{R}^d$ .
- Values  $v_i \in \mathbb{R}^d$  (say same dimension).
- All  $n$  queries use the same set of (key,value) pairs.

We effectively:

1. Build

$$S = A^\top A + \lambda I = \lambda I + \sum_{i=1}^n x_i x_i^\top$$

— cost  $O(nd^2)$ .

2. Build

$$B = A^\top Y$$

where  $Y$  stacks the value vectors — also  $O(nd^2)$ .

3. Invert the  $d \times d$  matrix  $S$ : cost  $O(d^3)$ .
4. Multiply  $S^{-1}B$ : cost  $O(d^3)$ .
5. Apply the resulting linear map to each of the  $n$  query vectors: each application is a matrix–vector multiply  $O(d^2)$ , so total  $O(nd^2)$ .

Given  $d < n$ , the term  $nd^2$  dominates  $d^3$  asymptotically, so the overall complexity is:

$$\boxed{O(nd^2)}.$$

So the correct option is  $O(nd^2)$ .

---

### (d) Backprop through ridge self-attention

We use ridge self-attention inside a Transformer, with a downstream loss.

The computation of the ridge solution is:

$$w^*(\lambda, \text{kevs. values}) \stackrel{\downarrow}{=} (A^\top A + \lambda I)^{-1} A^\top u.$$

The computation of the ridge solution is:

$$w^*(\lambda, \text{keys}, \text{values}) = (A^\top A + \lambda I)^{-1} A^\top y,$$

and then the attention outputs depend on  $w^*$  and the query vectors.

This is a composition of differentiable operations (matrix multiplications, inverse, etc.). Thus backprop can propagate gradients to:

- The **ridge**  $\lambda$  (if treated as a learnable parameter).
- The **keys** (they appear inside  $A$ ).
- The **values** (they appear inside  $y$  or  $Y$ ).
- The **queries** (used when evaluating the prediction on test inputs).

So **all four** items listed will receive gradients and can be learned.

---

#### (e) Cost of one Sherman–Morrison update

Sherman–Morrison:

$$(M + uv^\top)^{-1} = M^{-1} - \frac{1}{1 + v^\top M^{-1} u} (M^{-1} u)(v^\top M^{-1}).$$

Assuming we already know  $M^{-1}$ , and:

- Matrix–vector multiply costs  $O(d^2)$ ,
- Dyad  $uv^\top$  costs  $O(d^2)$ ,
- Inner product costs  $O(d)$ .

Steps:

1.  $z = M^{-1}u$ :  $O(d^2)$ .
2.  $w^\top = v^\top M^{-1}$  (or  $M^{-1}v$ ):  $O(d^2)$ .
3. Scalar  $\alpha = 1 + v^\top z$ :  $O(d)$ .
4. Dyad  $zw^\top$ :  $O(d^2)$ .
5. Scale and subtract from  $M^{-1}$ :  $O(d^2)$ .

Dominated by  $O(d^2)$ , so the correct choice is:

$O(d^2)$ .

#### (f) Efficient causal ridge self-attention

We want **causal** ridge self-attention with context length  $n$ :

- At time  $t$ , only keys/values from positions  $1, \dots, t$  are available and used.
- We want to avoid recomputing ridge regression from scratch at each  $t$  (which would be  $O(n^2 d^2)$ ).

Let  $k_t \in \mathbb{R}^d$  be the key,  $v_t \in \mathbb{R}^d$  be the value at time  $t$ .

Define:

- $S_t = \lambda I + \sum_{i=1}^t k_i k_i^\top \in \mathbb{R}^{d \times d}$ .
- $P_t = S_t^{-1}$  (the inverse).
- $B_t = \sum_{i=1}^t k_i v_i^\top \in \mathbb{R}^{d \times d}$ , so the ridge weight matrix at time  $t$  is  $W_t = P_t B_t$  (up to orientation).

**Initialization:**

- $S_0 = \lambda I$ ,  $P_0 = \lambda^{-1} I$ ,  $B_0 = 0$ .

**Recursive update at each time step  $t$ :**

1. **Update covariance:**

$$S_t = S_{t-1} + k_t k_t^\top.$$

2. **Update inverse using Sherman–Morrison** (rank-1 update with  $u = v = k_t$ ):

$$P_t = S_t^{-1} = P_{t-1} - \frac{P_{t-1} k_t k_t^\top P_{t-1}}{1 + k_t^\top P_{t-1} k_t}.$$

This is  $O(d^2)$ .

3. **Update cross-term  $B_t$ :**

$$B_t = B_{t-1} + k_t v_t^\top.$$

Again  $O(d^2)$  if values are  $d$ -dimensional.

4. **Compute output for query  $q_t$ :**

The multi-output ridge solution at time  $t$  is:

$$W_t = P_t B_t.$$

The prediction for query  $q_t$  is



okay for question 1c-onward, you did not consider the matrix multiplication time, can you factor that in and give me new answers for 1c, 1d, 1e, 1f

Thought for 3m 9s

Good catch — you're right that I'd implicitly treated matrix multiplies as "unit cost." Let's redo 1(c)-1(f) and explicitly include the cost of matrix/matrix–vector multiplications.

I'll use:

- $n$  = state dimension (so  $A \in \mathbb{R}^{n \times n}$ )
- $d$  = input dimension ( $u_k \in \mathbb{R}^d$ )
- $p$  = output dimension ( $y_k \in \mathbb{R}^p$ )
- Dense costs (standard):
  - matvec  $M \in \mathbb{R}^{a \times b}, x \in \mathbb{R}^b: \Theta(ab)$
  - matmat  $M_1 \in \mathbb{R}^{a \times b}, M_2 \in \mathbb{R}^{b \times c}: \Theta(abc)$

### 1(c) Efficient computation with convolutions (with matrix cost)

We assume the kernel  $K = \{K_\ell\}_{\ell=0}^L$  is already known, with  $K_\ell \in \mathbb{R}^{p \times d}$ , and we have input  $u_k \in \mathbb{R}^d$  of length  $L + 1$ .  
Output:

$$y_k = \sum_{\ell=0}^L K_\ell u_{k-\ell}.$$

#### Work and depth with convolutions

For each output time step  $k$ :

- We must compute up to  $L + 1$  **matvecs**  $K_\ell u_{k-\ell}$ , each costing  $\Theta(pd)$ .
- Then sum the resulting  $L + 1$  vectors in  $\mathbb{R}^p$ .

**Total work** (if you literally do the sum as written):

- Number of matvecs:  $\sim (L + 1)^2$  across all  $k \in \{0, \dots, L\}$ .
- Work:  $\Theta(L^2 \cdot pd)$ .

In practice you'd implement this as a convolution (FFT or highly optimized conv operator) with work closer to  $\Theta(Lpd \log L)$ , but the question is about parallelism / critical path, so:

#### Critical path length (with matmul cost):

- Suppose one matvec  $K_\ell u_j$  has latency  $T_{\text{mv}} = \Theta(pd)$  if done "straight-line."
- With enough cores:
  - All  $K_\ell u_{k-\ell}$  across all  $(k, \ell)$  can be done *in parallel* → contributes  $\sim T_{\text{mv}}$  to depth.
  - For a fixed  $k$ , summing  $L + 1$   $p$ -dimensional vectors via a tree reduction costs  $\Theta(\log L)$  additions per component, so depth  $\Theta(\log L)$ .

So the **minimum critical path length** for the convolution implementation is roughly

$$\text{depth}_{\text{conv}} \sim T_{\text{mv}} + \Theta(\log L) = \Theta(pd) + \Theta(\log L).$$

If you further parallelize within the matvec (over  $p$  and  $d$ ), the  $T_{\text{mv}}$  term itself shrinks toward  $\Theta(\log p + \log d)$ , so you get something like  $\Theta(\log p + \log d + \log L)$ .

### Naive direct computation from the recursion

Unrolled recursion:

$$x_{k+1} = Ax_k + Bu_k, \quad y_k = Cx_k + Du_k.$$

Per time step  $k$ :

- $Ax_k \in \mathbb{R}^n: \Theta(n^2)$  (dominant term).
- $Bu_k \in \mathbb{R}^n: \Theta(nd)$ .
- $Cx_k \in \mathbb{R}^p: \Theta(pn)$ .
- $Du_k \in \mathbb{R}^p: \Theta(pd)$ .

Total per-step work:  $\Theta(n^2 + nd + pn + pd)$ .

If  $d, p = O(n)$ , this is  $\Theta(n^2)$  per step, so

- **Total work:**  $\Theta(Ln^2)$ .
- **Critical path length:** each state depends on the previous, so you *must* do at least  $L$  sequential state

Total per-step work:  $\Theta(n^2 + nd + pn + pd)$ .

If  $d, p = O(n)$ , this is  $\Theta(n^2)$  per step, so

- **Total work:**  $\Theta(Ln^2)$ .
- **Critical path length:** each state depends on the previous, so you *must* do at least  $L$  sequential state updates. Even if each matvec is fully parallelized internally, you still have

$$\boxed{\text{depth}_{\text{recur}} \sim L \cdot T_{\text{step}}}$$

where  $T_{\text{step}}$  is the latency of one state update (essentially one  $n \times n$  matvec). If you parallelize inside the matvec, this is  $\Theta(L \log n)$ .

#### Summary of (c):

- Recurrent: work  $\Theta(Ln^2)$ , depth  $\Theta(L \cdot T_{\text{step}})$ .
- Convolution (with known kernel): work roughly  $\Theta(L^2pd)$  in the naive implementation, but depth only  $\Theta(T_{\text{mv}} + \log L)$ .  
So with enough parallel compute, the convolutional implementation has dramatically shorter critical path (factor  $\approx L/\log L$ ) at the cost of more total arithmetic.

### 1(d) Efficient kernel computation (with matrix cost)

We need to compute

$$K_0 = D, \quad K_\ell = CA^{\ell-1}B, \quad \ell = 1, \dots, L.$$

Here:

- $A \in \mathbb{R}^{n \times n}$
- $B \in \mathbb{R}^{n \times d}$
- $C \in \mathbb{R}^{p \times n}$ .

#### Naive sequential approach

One simple unrolled approach:

- Let  $M_1 = B$ , and define

$$M_{\ell+1} = AM_\ell \quad (\ell = 1, \dots, L-1).$$

Then  $M_\ell = A^{\ell-1}B$ .

- Then  $K_\ell = CM_\ell$ .

Costs:

- Each  $M_{\ell+1} = AM_\ell$ : matmat  $n \times n$  by  $n \times d$ :  $\Theta(n^2d)$ .
- Do this for  $\ell = 1, \dots, L-1 \rightarrow$  work  $\Theta(Ln^2d)$ .
- Each  $K_\ell = CM_\ell$ : matmat  $p \times n$  by  $n \times d$ :  $\Theta(pnd)$ , so total  $\Theta(Lpnd)$  (usually not the bottleneck if  $p$  small/ $\approx d$ ).

So naive **work**:  $\Theta(Ln^2d)$ .

**Depth** (sequential):  $\Theta(L)$  matmats, each of latency  $T_{\text{mm}} = \Theta(n^2d)$  (or  $\Theta(n^3)$  if you treat it as  $n \times n \times n \times n$ ), so:

$$\boxed{\text{depth}_{\text{naive}} \sim L \cdot T_{\text{mm}}}.$$

#### Parallel / divide-and-conquer strategy (for $L = 2^N$ )

Goal: reduce **depth** while keeping similar or slightly larger total work.

Two key ideas:

##### 1. Matrix-power precomputation via repeated squaring

Compute:

$$A, A^2, A^4, \dots, A^{2^{N-1}}$$

in  $\Theta(N) = \Theta(\log L)$  dense  $n \times n$  matmats.

Work:  $\Theta(n^3 \log L)$ .

Depth:  $\Theta(\log L) \cdot T_{n \times n}$ .

##### 2. Divide-and-conquer over time with parallelism

You can structure the computation of all  $M_\ell = A^{\ell-1}B$  as a segment tree or parallel prefix:

- At the top level, split the time indices  $\{1, \dots, L\}$  into first half and second half.
- Use the precomputed power  $A^{L/2}$  to "jump" from the end of the first half to the start of the second half.
- Recurse within each half; at each recursion you multiply by appropriate power of  $A$ .

This gives:

- **Work** still on the order of  $\Theta(Ln^2d)$  matmats  $\downarrow$  plus  $\Theta(n^3 \log L)$  for power precomputation.
- **Depth**: each level of recursion uses one additional multiplication by a precomputed power of  $A$ , so

Ask anything

+  Thinking



This gives:

- Work still on the order of  $\Theta(Ln^2d)$  matmats  $AM$  plus  $\Theta(n^3 \log L)$  for power precomputation.
- Depth: each level of recursion uses one additional multiplication by a precomputed power of  $A$ , so

$$\text{depth}_{\text{kernel}} \sim \log L \cdot T_{\text{mm}}$$

vs  $\sim L \cdot T_{\text{mm}}$  for the naive sequential power chain.

So factoring matrix-multiplication time:

- Naive kernel generation:  $\Theta(Ln^2d)$  work, depth  $\Theta(LT_{\text{mm}})$ .
- Divide-and-conquer / prefix-style: similar work (plus  $\Theta(n^3 \log L)$  overhead), but depth only  $\Theta(\log L \cdot T_{\text{mm}})$ .

### 1(e) Diagonal $A$ (with matrix cost)

Now suppose

$$A = \text{diag}(a_1, \dots, a_n).$$

Then:

- $A^{\ell-1} = \text{diag}(a_1^{\ell-1}, \dots, a_n^{\ell-1})$ ; powering is **elementwise**, no heavy matrix-matrix operations.

To compute  $K_\ell = CA^{\ell-1}B$ :

1. Compute  $A^{\ell-1}B$ . Because  $A^{\ell-1}$  is diagonal, this is just scaling each **row** of  $B$  by a scalar:
  - Cost per  $\ell$ :  $\Theta(nd)$  (rowwise scaling).
  - For all  $\ell$ :  $\Theta(Lnd)$ .
2. Then multiply by  $C$ : for each  $\ell$ ,  $K_\ell = C(A^{\ell-1}B)$  is a matmat  $p \times n$  by  $n \times d$ :  $\Theta(pnd)$ , so  $\Theta(Lpnd)$  overall.

You can improve further:

- Precompute the “base” rank-1 contributions:

$$M_r := C_{:,r}B_{r,:} \in \mathbb{R}^{p \times d}, \quad r = 1, \dots, n,$$

so:

$$K_\ell = \sum_{r=1}^n a_r^{\ell-1} M_r.$$

- Then:

- Precompute  $M_r$  once:  $\Theta(pnd)$ .
- Maintain the sequences  $a_r^{\ell-1}$  by recurrence  $a_r^\ell = a_r^{\ell-1}a_r$ :  $\Theta(Ln)$ .
- For each  $\ell$ , form  $K_\ell$  via a weighted sum of  $M_r$ :  $\Theta(npd)$  if done straightforwardly.

The main point: **multiplying by  $A$  is now cheap** — we never do an  $n \times n$  dense matmat, only scaling operations.

Compared to dense  $A$ :

- Dense: each  $AM$  is  $\Theta(n^2d)$ .
- Diagonal: each  $AM$  is  $\Theta(nd)$ .

So:

$$\boxed{\text{Kernel work with diagonal } A \sim \Theta(Lnd + Lpnd) \text{ vs } \Theta(Ln^2d) \text{ for dense } A.}$$

Depth-wise, you can still apply the same divide-and-conquer tricks as in (d), but the per-level cost is now  $\Theta(nd)$  instead of  $\Theta(n^2d)$ .

### 1(f) DPLR: $A = I_n + pp^\top$ (with matrix cost)

Now suppose

$$A = I_n + pp^\top, \quad p \in \mathbb{R}^n.$$

For any vector  $x \in \mathbb{R}^n$ ,

$$Ax = x + p(p^\top x).$$

This can be computed as:

1.  $s = p^\top x$ : inner product, cost  $\Theta(n)$ .
2.  $ps$ : scale  $p$  by scalar  $s$ , cost  $\Theta(n)$ .
3. Add to  $x$ : cost  $\Theta(n)$ .

So one **vector** multiplication by  $A$  costs  $\Theta(n)$  instead of  $\Theta(n^2)$ .

For a **matrix**  $M \in \mathbb{R}^{n \times d}$  (e.g., multiple columns or  $\downarrow$  iple starting vectors):

Ask anything

+ Thinking



So one **vector** multiplication by  $A$  costs  $\Theta(n)$  instead of  $\Theta(n^2)$ .

For a **matrix**  $M \in \mathbb{R}^{n \times d}$  (e.g., multiple columns or multiple starting vectors):

- Treat each column as a vector:
  - For each column  $m_j$ ,  $Am_j = m_j + p(p^\top m_j)$ .
  - Inner product per column:  $\Theta(n)$ .
  - Scaling and addition:  $\Theta(n)$ .
- Total cost:  $\Theta(nd)$  to compute  $AM$ , vs  $\Theta(n^2d)$  for dense  $A$ .

So for kernel generation:

- As in the naive scheme from (d), you can still compute

$$M_1 = B, \quad M_{\ell+1} = AM_\ell, \quad K_\ell = CM_\ell.$$

- But now each step  $M_{\ell+1} = AM_\ell$  is  $\Theta(nd)$ , not  $\Theta(n^2d)$ .

Hence:

Work with DPLR  $A \sim \Theta(Lnd + Lpnd)$  vs  $\Theta(Ln^2d)$  for dense  $A$ .

You can also apply the divide-and-conquer / prefix tricks:

- Use  $AM$  as the "combine" operator in a parallel scan over time.
- Each "combine" call costs  $\Theta(nd)$ .
- Total work still  $\Theta(Lnd)$  but depth becomes  $\Theta(\log L \cdot nd)$  if you treat each combine as a unit, or  $\Theta(\log L \cdot (\log n + \log d))$  if you parallelize inside the combines as well.

So the low-rank structure mainly **reduces the per-step cost** from  $\Theta(n^2d)$  to  $\Theta(nd)$ ; combining that with the same parallelization strategies as in (d) gives you both lower work and good depth.

If you want, we can plug in some concrete  $n, d, p, L$  (e.g., typical SSM numbers) and compare actual asymptotic regimes where recurrent vs convolution vs structured  $A$  changes which implementation is preferable.

□ ↗ ↙ ↪ ↤ ⋯



can you look at this solution and let me know if your answer for 4f and 4g matches

Thought for 23s

Short answer: yes, they match in substance. Let me line them up so you can see the correspondence.

[hw08\\_solution](#)

#### 4(f): Causal ridge self-attention

Official solution notation:

- Keys/values up to time  $t$ :  
 $A_t$  = stacked keys,  $B_t$  = stacked values
- Define  
 $M_t = \lambda I + A_t^\top A_t = \lambda I + \sum_{i=1}^t k_i k_i^\top$   
 $J_t = A_t^\top B_t = \sum_{i=1}^t k_i v_i^\top$
- Ridge weights at time  $t$ :  
 $W^*(t) = M_t^{-1} J_t$
- They then use **Sherman–Morrison** to update  $M_t^{-1}$  from  $M_{t-1}^{-1}$  in  $O(d^2)$ , and update  $J_t$  from  $J_{t-1}$  via a rank-1 update  $J_t = J_{t-1} + k_t v_t^\top$ .
- Output for query  $q_t$ : either multiply  $W^*(t)$  by  $q_t$ , or use the "minimal" variant  
 $y_t = W^*(t)^\top q_t = J_t^\top (M_t^{-1})^\top q_t$ , all in  $O(d^2)$  per time step.

What I wrote:

- I used almost the same objects, just different letters:
  - $S_t = \lambda I + \sum_{i=1}^t k_i k_i^\top$  (this is their  $M_t$ )
  - $P_t = S_t^{-1}$  (their  $M_t^{-1}$ )
  - $B_t = \sum_{i=1}^t k_i v_i^\top$  (their  $J_t$ )
- Recursion:
  - Update  $P_t$  from  $P_{t-1}$  using Sherman–Morrison on the rank-1 update  $k_t k_t^\top \rightarrow O(d^2)$ .
  - Update  $B_t = B_{t-1} + k_t v_t^\top$ . 
  - Weight matrix  $W_t = P_t B_t$ , and output  $y_t = W_t^\top q_t$ ; total per-step  $O(d^2)$ .

Ask anything

+  Thinking



- Recursion:
  - Update  $P_t$  from  $P_{t-1}$  using Sherman–Morrison on the rank-1 update  $k_t k_t^\top \rightarrow O(d^2)$ .
  - Update  $B_t = B_{t-1} + k_t v_t^\top$ .
  - Weight matrix  $W_t = P_t B_t$ , and output  $y_t = W_t^\top q_t$ ; total per-step  $O(d^2)$ .

So conceptually it's *exactly* the same algorithm:

- Track the inverse covariance  $M_t^{-1}$  and the "cross term"  $J_t$ ,
- Update them in  $O(d^2)$  per step via rank-1 updates,
- Compute each causal attention output in  $O(d^2)$ ,
- Total cost  $O(nd^2)$ .

The official solution goes further and writes the recursion for  $W^*(t)$  explicitly and then shows the "minimal" variant that only stores  $M_t^{-1}$  and  $J_t$ ; I described the same two ideas but in a slightly more high-level way.

There's no contradiction.

**Verdict for 4(f):** Same method and same complexity, just different notation and detail level.

#### 4(g): Attention weights for ridge-attention

**Official solution:**

They start from

$$W^* = (\lambda I + A^\top A)^{-1} A^\top B,$$

so for query  $q$ ,

$$W^{*\top} q = ((\lambda I + A^\top A)^{-1} A^\top B)^\top q = B^\top (A(\lambda I + A^\top A)^{-1}) q.$$

They interpret:

- Let  $w(q) := A(\lambda I + A^\top A)^{-1} q \in \mathbb{R}^n$ .
- Since  $B^\top = [v_1 \ v_2 \dots \ v_n]$ , the output is

$$\sum_{i=1}^n w_i(q) v_i,$$

so the **ridge-attention weights** are the entries of

$$w(q) = A(\lambda I + A^\top A)^{-1} q.$$

They then give the kernel-ridge version:

$$\text{weights} = (\lambda I_n + A A^\top)^{-1} A q.$$

**What I wrote:**

- I denoted the keys matrix as  $K \in \mathbb{R}^{n \times d}$  (their  $A$ ).
- I defined

$$P = (K^\top K + \lambda I)^{-1},$$

and

$$w(q) = K P q \in \mathbb{R}^n.$$

Then

$$\hat{y}(q) = \sum_{i=1}^n w_i(q) v_i.$$

If you substitute  $K = A$ ,  $P = (\lambda I + A^\top A)^{-1}$ , you get exactly

$$w(q) = A(\lambda I + A^\top A)^{-1} q,$$

which is the same expression they identify as the attention weights.

I didn't explicitly write the kernel-ridge alternative  $(\lambda I_n + A A^\top)^{-1} A q$ , but that's just the dual formulation of the same mapping.

**Verdict for 4(g):** Yes, my answer is the same: the weights are the  $n$ -vector

$$w(q) = A(\lambda I + A^\top A)^{-1} q,$$

whose entries tell you how each value vector is weighted in the final output.

If you'd like, I can rewrite my 4(f) / 4(g) in exactly their notation (with  $A_t, B_t, M_t, J_t$ ) so you can just drop them straight into your notes.

Sources



↶ ↷ ↸ ↹ ⋮

Ask anything

+  Thinking



